

Joint Bidding and Geographical Load Balancing for Datacenters: Is Uncertainty a Blessing or a Curse?

Ying Zhang*, Lei Deng*, Minghua Chen*, and Peijian Wang†

*Department of Information Engineering, The Chinese University of Hong Kong, Hong Kong

†School of Electrical and Information Engineering, Xi'an Jiaotong University, Xi'an, China

Abstract—We consider the scenario where a cloud service provider (CSP) operates multiple geo-distributed datacenters to provide Internet-scale service. Our objective is to minimize the total electricity and bandwidth cost by jointly optimizing electricity procurement from wholesale markets and geographical load balancing (GLB), *i.e.*, dynamically routing workloads to locations with cheaper electricity. Under the ideal setting where exact values of market prices and workloads are given, this problem reduces to a simple LP and is easy to solve. However, under the realistic setting where only distributions of these variables are available, the problem unfolds into a non-convex infinite-dimensional one and is challenging to solve. Our main contribution is to develop an algorithm that is proven to solve the challenging problem optimally and efficiently, by exploring the full design space of strategic bidding. Trace-driven evaluations corroborate our theoretical results, demonstrate fast convergence of our algorithm, and show that it can reduce the cost for the CSP by up to 20% as compared to baseline alternatives. Our study highlights the intriguing role of uncertainty. While variability in workloads deteriorates the cost-saving performance of joint electricity procurement and GLB, counter-intuitively, variability in market prices can be exploited to achieve a cost reduction even *larger* than the setting without price variability.

I. INTRODUCTION

As cloud computing services become prevalent, the electricity cost of worldwide datacenters hosting these services has skyrocketed, reaching \$16B in 2010 [19]. Electricity cost represents a large fraction of the datacenter operating expense [33], and it is increasing at an alarming rate of 12% annually [4]. Consequently, reducing electricity cost has become a critical concern for datacenter operators [27].

There have been substantial research on reducing power consumption and related cost of datacenters [14], [16], [31], [32]. Among them, geographical load balancing (GLB) is a promising technique [27], [28], [30]. By *dynamically* routing workloads to locations with cheaper electricity, GLB has been shown to be effective in reducing electricity cost (*e.g.*, by 2–13% [27]) of geo-distributed datacenters operated by a CSP. Many existing works explore price *diversity across geographical locations* to reduce electricity cost [27], [28], [35]. Some recent studies also advocate additional price *diversity across time* at a location, by for example using electricity storage system and demand response for arbitrage [31] or opportunistically optimizing various electricity procurement options [6], [13], [34].

Inspired by these advances and recent practices that CSPs moving into electricity markets (*e.g.*, Google Energy LLC [15]), we consider the scenario where a CSP jointly performing GLB and electricity procurement from deregulated markets. The market prices are set by running *auction* mechanisms




Fig. 1. (a) We fix market prices to their means and increase standard deviations of workloads. Cost reductions of our solution and baseline decrease as the standard deviations increase. (b) We fix workloads to their means and increase standard deviations of prices. Cost reduction of our solution increases as the standard deviations increase, while that of baseline stays constant. More details are in our technical report [38].

among the electricity suppliers and consumers, *cf.*, [37]. The goal is to minimize the total electricity and bandwidth cost, by exploiting price diversity in both geographical locations (by GLB) and time (by procurement in local sequential markets).

Under the ideal setting where exact values of market prices and workloads are given, the problem reduces to a simple LP and is easy to solve, by for an example solution in [27]. In practice, however, the actual values of these variables are revealed only at the operating time, and only their distributions are available when procuring electricity by submitting bids to markets (bidding). Under such realistic settings, the problem unfolds into a non-convex infinite-dimensional one. Our focus in this paper is to develop an algorithm to solve the problem optimally.

Our study highlights the intriguing role of uncertainty. On one hand, workload variability undermines the efficiency of balancing supply and demand (proportional to workload) on electricity markets. As a result, the cost-saving performance of joint bidding and GLB deteriorates as workload variability increases, as illustrated in Fig. 1(a). On the other hand, counter-intuitively, higher variability in market prices allows us to extract larger *coordination* gain in sequential procurement in day-ahead and real-time markets [3], [12], [22]. As shown in Fig. 1(b), capitalizing such gain leads to a cost reduction even *larger* than the setting without price variability.

In our solution, we explore the full design space of strategic bidding to *simultaneously* exploit the price variability and combat the workload variability, so as to maximize the cost saving. We make the following contributions.

▷ We present necessary backgrounds on electricity markets in Sec. II. Then in Sec. III, we formulate the problem of cost minimization by joint bidding and GLB, under the realistic setting where only distributions of market prices and

workloads are available. The problem is a non-convex infinite-dimensional one and is in general challenging to solve.

▷ To address the non-convexity challenge, in Sec. IV, we leverage problem structures to characterize a subregion of the feasible set so that (i) it contains the optimal solution, and (ii) the problem over this subregion becomes a convex one. We then solve the reduced problem by a nested-loop solution.

▷ In the inner loop, we fix the GLB decision and optimize bidding strategies for local sequential markets. We derive an easy-to-compute closed-form optimal solution in Sec. IV-B. The optimal bidding strategies not only address the infinite-dimension challenge, but also allow the CSP to simultaneously exploit price variability and combat workload variability.

▷ In the outer loop, we solve the remaining GLB problem given optimal bidding strategies. While the problem is convex and of finite dimension, its objective function does not admit an explicit-form expression. Consequently, its gradient cannot be computed explicitly, and gradient/subgradient-based algorithms cannot be directly applied. In Sec. IV-C, we tackle this issue by adapting a zero-order optimization algorithm, named General Pattern Search (GPS) [20], to solve the problem without knowing the explicit-form expression of the objective function. Finally, we prove that our nested-loop algorithm solves the joint bidding and GLB problem optimally. We discuss the computational complexity and issues related to practical implementation in our technical report [38].

▷ By evaluations based on real-world traces in Sec. VI, we show that our solution converges fast and reduces the CSP cost by up to 20% as compared to baseline alternatives.

Our study also adds understanding to energy cost management for entities other than datacenters. For example, [26] and [22] considered similar problems for utilities and microgrids, without fully exploring the bidding design space or pursuing optimal solution. Results of our study thus can help to optimize bidding strategy designs under such settings.

We discuss related works in Sec. VII and conclude the paper in Sec. VIII. Due to the space limitation, all proofs are deferred to [38].

II. ELECTRICITY MARKET PRELIMINARY

In a region, there are two electricity wholesale markets, *day-ahead* market and *real-time* market, to balance the electricity supply and demand in two timescales. We show the critical operations in Fig. 2 and explain the details in the following.




Fig. 2. Operation of day-ahead market and real-time market.

Day-Ahead Market. The day-ahead market is a forward market to trade the electricity one day before dispatching. The electricity supply is auctioned in the day-ahead market. The sellers, *i.e.*, generation companies, submit (hourly) generation offers, and the buyers, *i.e.*, utilities or CSPs, submit (hourly)

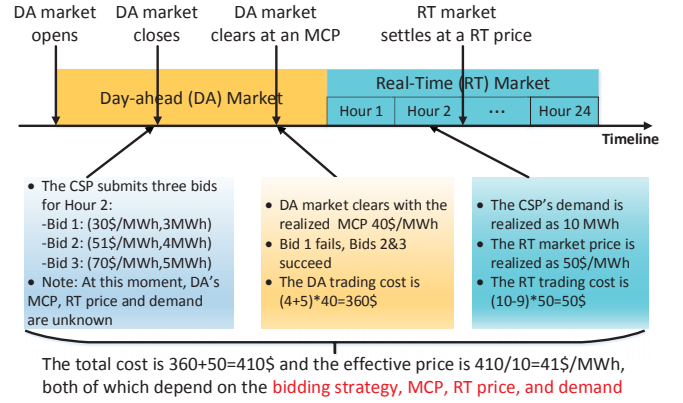


Fig. 3. An illustrating example for the CSP to participate in markets.

demand bids, all in the format of $\langle \text{marginal price}, \text{quantity} \rangle$, to the *auctioneer*, *i.e.*, the Independent System Operator (ISO).

In the offers (resp. bids), the generation companies (resp. utilities and CSPs) specify the amount of electricity they want to sell (resp. buy) and at which marginal price. Each seller (resp. buyer) is allowed to submit *multiple* offers (resp. bids) [3] in the same auction with different prices and quantities. The ISO matches the offers with the bids, typically using a well-established double auction mechanism [37]. The outcome of the auction is that it determines a *market clearing price* (MCP) for all the traded units. The bids with prices higher than MCP and the offers with prices lower than MCP will be accepted, and the electricity will be traded at MCP. Upon day-ahead market settlement, the generation companies (resp. utilities and CSPs) will be notified the quantity and MCP of electricity that they commit to generate (resp. consume).

The actual value of MCP is revealed only after the day-ahead market is settled/cleared, and they are unknown to market participants at the time of submitting bids/offers.

We show an example in Fig. 3 from the perspective of our CSP. Suppose that the CSP submits three bids to the day-ahead market: $\langle 30\$/\text{MWh}, 3\text{MWh} \rangle$, $\langle 51\$/\text{MWh}, 4\text{MWh} \rangle$, $\langle 70\$/\text{MWh}, 5\text{MWh} \rangle$. Now if ISO announces that the MCP is $40\$/\text{MWh}$ after the auction, then the second and the third bid will be accepted since their bidding prices are higher than MCP. Thus the CSP gets $4 + 5 = 9\text{MWh}$ of day-ahead committed supply at the price of MCP, *i.e.*, $40\$/\text{MWh}$. The day-ahead trading cost is thus $9 \times 40 = 360\$$.

Real-Time Market. The mismatch between day-ahead committed supply (as discussed above) and real-time demand is balanced on the real-time market, in a pay-as-you-go fashion. In particular, the system calls the short-start fast-responding generating units, which is usually more expensive, to standby and meet the instantaneous power shortage if any. The real-time price is set after the real-time dispatching and are not exactly known a priori.

- In case that the day-ahead committed supply matches exactly the actual demand, there is no real-time cost.
- In case of under-supply, (*i.e.*, the committed supply is less than the real-time demand), the CSP will pay for extra supply at the real-time price.
- In case of over-supply, the system needs to reduce the

power generation output or pay to schedule elastic load [24] to balance the supply, both incurring operational overhead and consequently economic loss. In this case, the CSP will receive a rebate at price $\beta \cdot \text{MCP}$ for the unused electricity (recall that the planned supply is purchased from the day-ahead market at price MCP). Here $\beta \in [0, 1)$ is a discounting factor capturing the overhead-induced cost in handling over-supply situation.

The overall electricity cost for the CSP is the sum of day-ahead procurement cost and the real-time settlement cost, which can be in the form of extra payment or rebate. A more concrete market example can be found in [17]. We remark that the framework developed in this paper can also be extended to different pricing models described in [12], [24], [26] (see our discussions in [38]).

Back to our example for the CSP in Fig. 3, suppose that the CSP's real-time demand is 10MWh. Since the day-ahead committed supply is only 9MWh, *i.e.*, the under-supply case happens, the CSP needs to buy 1MWh extra electricity from the real-time market. Now if the real time price is 50\$/MWh, the real-time trading cost of the CSP will be $1 \times 50 = 50\$$. The total cost is the sum of day-ahead trading cost and real-time trading cost, which is $360+50=410\$$.

Cost Structure. An important observation is that the overall cost depends on not only the actual demand, the day-ahead MCP and the real-time price, but also the mismatch between the day-ahead committed supply and the actual demand. As the day-ahead committed supply depends on day-ahead market bidding strategy of the CSP, the overall cost is thus also a function of the bidding strategy. We remark that such cost structure is unique to electricity procurement in electricity markets and motivates the bidding strategy design [26].

III. SYSTEM MODEL AND PROBLEM FORMULATION

We consider the scenario of a CSP providing computing-intensive services (*e.g.*, Internet search) to users in N regions by operating N geo-distributed datacenters, one in each region, as exemplified in Fig. 4. Service workloads from a region can be served either by the local datacenter or possibly by datacenters in other regions through GLB. The CSP directly participates in wholesale electricity markets in each region, to obtain electricity to serve the local datacenter. Based on (i) distributions of hourly service workloads and (ii) distributions of market settlement prices, the CSP aims at minimizing the expected total operating cost by optimizing GLB and bidding strategies on the markets. The hourly timescale aligns with both the settlement timescale in wholesale markets [30] and the suggested time granularity for performing GLB [27].

Without loss of generality, we focus on minimizing cost of a particular operation hour of the CSP, as shown in Fig. 2.

A. Workload and Geographical Load Balancing

Workload and Electricity Demand. We assume that each datacenter is power-proportional, which means that its electricity consumption is proportional to its workload [27]. For example, Google reports that each search requires about 0.28Wh electricity for its datacenters [27]. Without loss of generality, we assume that the workload-to-electricity coefficient is unit




Fig. 4. The scenario that we consider in this paper.

one for all datacenters and thus use the workload served by a datacenter to represent its electricity demand. Our results can be easily generalized to the case where the coefficients are different for different datacenters.

We model the workload originated from region i as a random variable U_i in the range $[\underline{u}_i, \bar{u}_i]$, with a probability density function (PDF) $f_{U_i}(u)$ that can be empirically estimated from historical data. We assume that all U_i 's are independent.

Geographical Load Balancing.¹ We denote the GLB decision by $\alpha = [\alpha_{ij} : i, j = 1, \dots, N] \in \mathbb{R}_{N \times N}$ which satisfies

$$\sum_j \alpha_{ij} \geq 1, \quad \forall i = 1, \dots, N, \quad (1)$$

$$\alpha_{ii} \geq \lambda_i, \quad \forall i = 1, \dots, N, \quad (2)$$

$$\bar{v}_j \triangleq \sum_{i=1}^N \alpha_{ij} \bar{u}_i \leq C_j, \quad \forall j = 1, \dots, N. \quad (3)$$

$$0 \leq \alpha_{ij} \leq 1, \quad \forall i, j = 1, \dots, N, \quad (4)$$

$$\alpha_{ij} = 0, \quad \forall (i, j) \in \mathcal{G}, \quad (5)$$

where $\mathcal{G} \triangleq \{(i, j) \mid \text{workloads from region } i \text{ cannot be routed to datacenter } j\}$ captures the topological constraints.

Here α_{ij} represents the fraction of the workload originated from region i that will be routed to datacenter j . Constraints in (1) mean that all workloads must be served. Constraints in (2) capture that λ_i fraction of the workload originated from region i can only be served locally due to various reasons such as delay requirements. Constraints in (3) ensure that the total workload coming into datacenter j can be served even in the largest realization of workload. Constraints in (5) describe that the workload cannot be routed to a datacenter that is too far away from its own region. We define the set of all feasible GLB decisions as

$$\mathcal{A} \triangleq \{\alpha \in \mathbb{R}^{N \times N} \mid \alpha \text{ satisfies (1) - (5)}\}. \quad (6)$$

¹Under the conventional setting where datacenters obtain electricity from utilities, GLB is performed in CSP's real-time operation. Under the setting we consider in this paper, CSP needs to bid for electricity in the day-ahead market, where the amount of electricity to bid is a function of GLB decisions. As such, we consider doing joint GLB and electricity bidding in CSP's day-ahead operation, in order to fully explore the new design space enabled by the setting considered in this paper. It is conceivable to perform GLB in both day-ahead and real-time operations of CSP to further minimize the energy cost, which we leave for future study.

Given the GLB decision α , the total workload for datacenter j is given by $V_j = \sum_i \alpha_{ij} U_i$. Since $U_i, \forall i$ are random variables, V_j is also a random variable with a PDF

$$f_{V_j}(v) = f_{U_{1j}} \otimes f_{U_{2j}} \otimes \dots \otimes f_{U_{Nj}}(v), \quad (7)$$

where \otimes is the convolution operator and the distribution functions in the convolution are given by

$$f_{U_{ij}}(u) = \begin{cases} \frac{1}{\alpha_{ij}} f_{U_i} \left(\frac{u}{\alpha_{ij}} \right), & \text{if } \alpha_{ij} > 0, \\ \delta(u), & \text{if } \alpha_{ij} = 0, \end{cases} \quad (8)$$

where $\delta(\cdot)$ denotes Dirac delta function.

Bandwidth Cost. According to our investigation, the bandwidth cost can be comparable to the electricity cost [38]. Let $z_{ij} \geq 0$ be the unit bandwidth cost from region i to datacenter j . The expected network cost of routing the workload to different datacenters is given by

$$\text{BCost}(\alpha) = \sum_{i=1}^N \sum_{j=1}^N z_{ij} \cdot \alpha_{ij} \cdot \mathbb{E}(U_i). \quad (9)$$

B. Electricity Market Price and Bidding Curve

Day-ahead MCP and Real-time Market Price. At the time of making joint bidding and GLB decisions, MCPs of day-ahead markets in N regions are unknown. We model them as N independent random variables P_j ($j \in [1, N]$), each with probability distribution $f_{P_j}(p)$ that can be empirically estimated from historical data [6]. Here we assume that the CSP has negligible market power and its bidding and GLB behavior will not affect the dynamics of electricity markets².

Similarly, the real-time market prices in N regions are also unknown when making bidding and GLB decisions. We model the price of real-time market j as a random variable P_j^{RT} whose probability distribution can also be empirically estimated from historical data [6]. We define $\mu_j^{\text{RT}} \triangleq \mathbb{E}[P_j^{\text{RT}}]$ as the expectation of P_j^{RT} . We assume that all day-ahead MCPs P_j 's and real-time market prices P_j^{RT} 's are independent³.

Bidding Curve. We explore the full design space of bidding strategy via *bidding curve*, which is a well-accepted concept in the power system community [12], [22]. Bidding curve, denoted as $q_j(p)$, is a function that maps the (realized) day-ahead market MCP to the amount of electricity the CSP wishes to obtain from day-ahead market j , by placing multiple bids. We remark that it is a common practice for one entity (*e.g.*, a utility company) to submit multiple bids to one electricity market.

Bidding curve is useful in designing bidding strategies in the following sense. First, any set of bids can be mapped to a bidding curve. Suppose the CSP submits K bids, namely $\langle b_j^k, q_j^k \rangle, k = 1, \dots, K$, to the day-ahead market of region j , where b_j^k is the bidding price and q_j^k is the bidding




Fig. 5. An illustrating example for the (step-wise) bidding curve constructed from the submitted three bids in Fig. 3.

quantity of the k -th bid. The corresponding bidding curve is a step-wise decreasing function as

$$q_j(p) = \sum_{k: b_j^k \geq p} q_j^k, \quad \forall p \in \mathbb{R}^+. \quad (10)$$

For example, considering the three bids in Fig. 3, we can construct the corresponding bidding curve as shown in Fig. 5.

Recall that if day-ahead market MCP is p , then all bids whose bidding prices are higher than p will be accepted. Thus, the right hand side of (10) represents the total amount of electricity obtained when the day-ahead MCP is p . Clearly, the purchased amount will be non-increasing in MCP p . Thus, a valid bidding curve $q_j(p)$ must be a non-increasing function.

Second, any non-increasing function is a valid bidding curve and can be realized by placing a set of bids. For example, the bidding curve in (10) can be realized by placing the K bids $\langle b_j^k, q_j^k \rangle, k = 1, \dots, K$ stated above.

Based on the above two observations, we design bidding strategy by choosing a bidding curve from the feasible set

$$\mathcal{Q} \triangleq \{q(p) \mid q(p_1) \leq q(p_2), \forall p_1 \geq p_2, p_1, p_2 \in \mathbb{R}^+\}. \quad (11)$$

Remark. In this paper, we assume that the CSP is allowed to submit any number, possibly infinite number, of bids. This assumption allows us to significantly simplify the derivation of optimal solution to the joint bidding and GLB problem in Sec. IV. In Sec. V, we discuss how to approximately realize a continuous bidding curve with a limited number of bids in the practical implementation. Our simulation results in Sec. VI (Tab. I) suggest that the performance loss due to the approximation error is minor.

Electricity Cost. Given the bidding curve $q_j(p)$ and the GLB decision α , we denote the *expected* electricity procurement cost of the CSP in electricity market j as $\text{ECost}_j(q_j(p), \alpha)$, which consists of settlement in both day-ahead trading and real-time trading.

- In day-ahead trading, suppose that the MCP in the day-ahead market j is p , the committed supply will be $q_j(p)$ and the day-ahead trading cost is $p \cdot q_j(p)$.
- In real-time trading, the day-ahead committed supply $q_j(p)$ may not exactly match the real-time demand V_j . If $V_j = v$ and $v > q_j(p)$, *under-supply* happens and we need to buy $v - q_j(p)$ amount of electricity at expected price μ_j^{RT} , so the expected cost due to under-supply would be $\mu_j^{\text{RT}} \int_{q_j(p)}^{V_j} (v - q_j(p)) f_{V_j}(v) dv$. Similarly, if *over-supply* happens, the unused electricity $(q_j(p) - v)$ will be sold back at a discounted price βp and the expected rebate due to over-supply is $\beta p \int_0^{q_j(p)} (q_j(p) - v) f_{V_j}(v) dv$. The

²The assumption is reasonable as, *e.g.*, datacenters in the US only consume 2% of total electricity [11], and it is usually used in the literature such as [30].

³We remark that this independence assumption may not hold in practice. But it significantly simplifies our analysis and allows us to reveal some important insights. A comprehensive study of considering correlations between day-ahead MCPs and real-time prices would be an interesting future work.

$$\text{ECost}_j(q_j(p), \alpha) = \underbrace{\int_0^{+\infty} f_{P_j}(p) [pq_j(p)]}_{\text{Day-ahead trading cost}} - \underbrace{\beta p \int_0^{q_j(p)} (q_j(p) - v) f_{V_j}(v) dv}_{\text{Rebate of over-supply}} + \underbrace{\mu_j^{\text{RT}} \int_{q_j(p)}^{\bar{v}_j} (v - q_j(p)) f_{V_j}(v) dv}_{\text{Cost of under-supply}} dp. \quad (12)$$

Real-time trading cost

expected real-time trading cost is simply the under-supply cost minus the over-supply rebate (see Eq. (12)).

Based on the above analysis, we obtain the expression of $\text{ECost}_j(q_j(p), \alpha)$ in (12) by applying the total expectation theorem. Note that $\text{ECost}_j(q_j(p), \alpha)$ is related to the GLB decision α through the distribution of V_j (the workload of datacenter j), which is computed by (7) and (8).

C. Problem Formulation

We now formulate the problem of joint bidding and GLB:

$$\begin{aligned} \mathbf{P1}: \quad & \min \sum_{j=1}^N \text{ECost}_j(q_j(p), \alpha) + \text{BCost}(\alpha) \\ \text{var.} \quad & \alpha \in \mathcal{A}, q_j(p) \in \mathcal{Q}, j = 1, \dots, N. \end{aligned}$$

where \mathcal{A} is the set of all feasible GLB decisions, defined in (6) and \mathcal{Q} is the set of all feasible bidding curves, defined in (11). It is straightforward to see both \mathcal{A} and \mathcal{Q} are convex sets. The objective is to minimize the summation of electricity cost of N datacenters and network cost, by optimizing bidding strategies and GLB decisions. The consideration of joint bidding and GLB as well as the market and demand uncertainty differentiates our work from existing works, e.g., [6], [27], [28], [33]. We emphasize that it is important to consider input uncertainty to fully capitalize the economic benefit of joint bidding and GLB under real-world market mechanisms.

Challenges. There are two challenges in solving problem **P1**. First, it can be shown that the objective function of **P1** is non-convex with respect to $q_j(p)$ (see our technical report [38]). Second, the optimization variable $q_j(p)$ is a functional variable with infinite dimensions. Thus it is highly non-trivial to solve this non-convex infinite-dimensional problem optimally without incurring forbidden complexity.

IV. AN OPTIMAL JOINT BIDDING AND GLB SOLUTION

In this section, we design an algorithm to solve the challenging problem **P1** optimally and efficiently.

A. Reducing P1 to a Convex Problem and Approach Sketch

To begin with, we define a sub-region of \mathcal{Q} as follows

$$\hat{\mathcal{Q}}_j = \{q_j(p) | q_j(p) \in \mathcal{Q}, \text{ and } q_j(p) = 0, \forall p \geq \mu_j^{\text{RT}}\}. \quad (13)$$

As compared to \mathcal{Q} defined in (11), the new constraint in the definition of $\hat{\mathcal{Q}}_j$, i.e., $q_j(p) = 0, \forall p \geq \mu_j^{\text{RT}}$, means that we do not submit any bid to day-ahead market j with bidding price higher than μ_j^{RT} , i.e., the expected price of real-time market j . It is easy to verify that both \mathcal{Q} and $\hat{\mathcal{Q}}_j$ are convex sets.

Theorem 1: The following problem **P2** is *convex* and has the same optimal solution as **P1**:

$$\begin{aligned} \mathbf{P2}: \quad & \min \sum_{j=1}^N \text{ECost}_j(q_j(p), \alpha) + \text{BCost}(\alpha) \\ \text{var.} \quad & \alpha \in \mathcal{A}, q_j(p) \in \hat{\mathcal{Q}}_j, j = 1, \dots, N. \end{aligned}$$

Remarks. (i) Problems **P1** and **P2** differ only in the feasible set of bidding curve $q_j(p)$. It is \mathcal{Q} in **P1** but $\hat{\mathcal{Q}}_j$ in **P2**. The objective function is nonconvex over \mathcal{Q} but convex over $\hat{\mathcal{Q}}_j$, as shown in the proof of Theorem 1 in [38]; hence, **P1** is a nonconvex problem but **P2** now is a convex one. (ii) Intuitively, the optimal bidding curve for day-ahead market j must be in $\hat{\mathcal{Q}}_j$. This is because the CSP can always buy electricity from real-time market j at an expected price μ_j^{RT} ; thus it is not economic to submit bids with bidding price higher than μ_j^{RT} to day-ahead market j . Such bidding strategies must be in set $\hat{\mathcal{Q}}_j$, defined in (13).

Theorem 1 allows us to solve **P1** by solving the convex problem **P2**. However, **P2** still suffers the infinite-dimension challenge, since optimizing bidding curves in general requires us to specify the value of $q_j(p)$ for every $p \in [0, \mu_j^{\text{RT}})$. To proceed with our design, we first rewrite problem **P2**,

$$\begin{aligned} & \min_{\alpha \in \mathcal{A}} \min_{q_j(p) \in \hat{\mathcal{Q}}_j, \forall j} \left\{ \sum_{j=1}^N \text{ECost}_j(q_j(p), \alpha) + \text{BCost}(\alpha) \right\} \\ & = \min_{\alpha \in \mathcal{A}} \left\{ \sum_{j=1}^N \underbrace{\left[\min_{q_j(p) \in \hat{\mathcal{Q}}_j} \text{ECost}_j(q_j(p), \alpha) \right]}_{\text{Problem EP}_j(\alpha), \text{ solved in Sec. IV-B}} + \text{BCost}(\alpha) \right\} \quad (14) \\ & \quad \text{Problem P3, solved in Sec. IV-C} \end{aligned}$$

The structure of the expression in (14) suggests a nested-loop approach to solve problem **P2**.

- **Inner Loop:** The CSP optimizes its bidding strategies for each regional day-ahead market with given GLB decision α , by solving the following problems:

$$\mathbf{EP}_j(\alpha) : \min_{q_j(p) \in \hat{\mathcal{Q}}_j} \text{ECost}_j(q_j(p), \alpha), \quad j = 1, \dots, N.$$

- **Outer Loop:** After solving the inner-loop problems $\mathbf{EP}_j(\alpha)$ and obtaining the optimal bidding curves, denoted by $q_j^*(p; \alpha), \forall j = 1, \dots, N$, the CSP optimizes the (finite-dimensional) GLB decision α by solving the following problem:

$$\mathbf{P3}: \min_{\alpha \in \mathcal{A}} \sum_{j=1}^N \text{ECost}_j(q_j^*(p; \alpha), \alpha) + \text{BCost}(\alpha).$$

In our technical report [38], we prove that the inner-loop problem $\mathbf{EP}_j(\alpha)$ and outer-loop problem **P3** are both convex, which are perhaps not surprising. In the next two subsections, we solve $\mathbf{EP}_j(\alpha)$ and **P3** to obtain an optimal joint bidding and GLB solution to **P2**, which is also optimal for **P1**.

B. Inner Loop: Optimal Bidding Given GLB Decision

The inner-loop problem $\mathbf{EP}_j(\alpha)$ is concerned about designing optimal bidding strategy for day-ahead market in region

j (by choosing $q_j(p) \in \hat{Q}_j$) with GLB decision α given, in face of demand and price uncertainty. Note that $\mathbf{EP}_j(\alpha)$ is closely related to the classic Newsvendor problem [18]. In the Newsvendor problem, the market prices are given and only the buying quantity should be optimized under demand uncertainty while we need to optimize both the bidding quantities and bidding prices simultaneously under both price and demand uncertainties.

Let the cumulative distribution function (CDF) of V_j , i.e., the demand of datacenter j , be $F_{V_j}(x) \triangleq \int_0^x f_{V_j}(v)dv$, where $f_{V_j}(v)$ is PDF of V_j given in (7). The following theorem shows that $\mathbf{EP}_j(\alpha)$ admits a closed-form solution $q_j^*(p; \alpha)$, addressing the infinite-dimension challenge.

Theorem 2: Given GLB decision α , we assume that $F_{V_j}(x)$ is strictly increasing; thus its inverse exists and is denoted as $F_{V_j}^{-1}(x)$. The optimal bidding curve for solving $\mathbf{EP}_j(\alpha)$ is given by, for $j = 1, \dots, N$,

$$q_j^*(p; \alpha) = \begin{cases} F_{V_j}^{-1}\left(\frac{\mu_j^{\text{RT}} - p}{\mu_j^{\text{RT}} - \beta p}\right), & \text{if } p \in [0, \mu_j^{\text{RT}}); \\ 0, & \text{otherwise.} \end{cases} \quad (15)$$

The optimal bidding curve $q_j^*(p; \alpha)$ is *universal* in that it does not depend on the distribution of day-ahead MCP P_j . This is because $q_j^*(p; \alpha)$ actually minimizes the expected electricity procurement cost for any p . It is easy to extend this result to the general non-decreasing $F_{V_j}(x)$ [38].

Robust to workload/demand variability. The optimal bidding curve $q_j^*(p; \alpha)$ in (15) is a decreasing function in the day-ahead MCP p . This meets our intuition that a good bidding strategy should purchase more electricity when p is low and less electricity when p is high. Similarly, it also meets our intuition that $q_j^*(p; \alpha)$ increases in the self-back discounting factor β . A larger β means less penalty of over-supply, thus the CSP should bid for more electricity from the day-ahead market to balance the penalty of under-supply. Bidding curves with these two features allow our overall solution to be robust to workload/demand variability, as shown in our case study in Fig. 1(a) as well as simulation results in Sec. VI-D.

Opportunistic to MCP variability. The optimal bidding curve $q_j^*(p; \alpha)$ in (15) also allows us to leverage variability in MCP of day-ahead market j to save cost. In particular, the unique two-sequential-market structure where the real-time market serves as a backup for the day-ahead market allows our bidding strategy $q_j^*(p; \alpha)$ to fully explore the benefit of low MCP values but control the risk of high MCP values. We elaborate as follows. When MCP fluctuates, its value, denoted by p , takes small and large values. When p is small, we can purchase cheap electricity from the day-ahead market and thus enjoys “gain”. When p is large, we have to purchase expensive electricity from the day-ahead market and thus suffers “loss”. However, when $p \geq \mu_j^{\text{RT}}$, our optimal bidding strategy $q_j^*(p; \alpha)$ will not purchase any electricity from the day-ahead market but purchase all electricity from the real-time market at the expected price μ_j^{RT} , bounding the “loss” due to high MCP values. Overall, the gain out-weights the loss and we achieve cost saving by leveraging MCP variability. In fact, the larger the MCP variability, the more significant the saving, as illustrated in our case study in Fig. 1(b).

Algorithm 1 An Algorithm for Solving **P3** Optimally

```

1: initialize  $\alpha^0 \leftarrow \mathbf{I}_{N \times N}$ ,  $t \leftarrow 0$ 
2: while not converge do
3:   current_value  $\leftarrow \mathbf{P3-Obj}(\alpha^t)$ 
4:   Get  $\alpha^{t+1}$  by invoking P3-Obj and comparing with
     current_value at most  $2N^2$  times (see [20, Fig. 3.4])
5:    $t \leftarrow t + 1$ 
6: end while
7:  $\alpha^* \leftarrow \alpha^t$ 
8: Compute  $q_j^*(p; \alpha^*)$  by (15) for all  $j \in [1, N]$ 
9: return  $\alpha^*$ ,  $q_j^*(p; \alpha^*)$  for all  $j \in [1, N]$ 
    
```

A subroutine to compute the objective value of P3

```

10: function P3-Obj( $\alpha$ )
11:   initialize  $j \leftarrow 1$ , val  $\leftarrow \text{BCost}(\alpha)$  by (9)
12:   while  $j \leq N$  do
13:     Compute  $q_j^*(p; \alpha)$  by (15)
14:     val  $\leftarrow \text{val} + \text{ECost}_j(q_j^*(p; \alpha), \alpha)$  by (12)
15:      $j \leftarrow j + 1$ 
16:   end while
17:   return val
18: end function
    
```

C. Outer Loop: Optimal GLB with Optimal Bidding Curve as a Function of GLB Decision

After obtaining the optimal bidding strategy $q_j^*(p; \alpha)$ as a function of GLB decision α , we now solve the outer-loop problem **P3** for optimizing GLB. While **P3** is convex and of finite dimension, its objective function does not admit an explicit-form expression since we do not have an explicit expression of the optimal objective value of $\mathbf{EP}_j(\alpha)$. Thus, gradient-based algorithms cannot be directly applied.

We tackle this issue by adapting a zero-order optimization algorithm, named General Pattern Search (GPS) [20], to solve the out-loop problem without knowing explicit expression of the objective function. *Zero-order optimization algorithms* are widely used to solve optimization problems without directly accessing the derivative information. The GPS algorithm in [20] is a popular zero-order optimization algorithm for solving problems with linear constraints, which is suitable for **P3**.

Our adapted GPS algorithm is an iterative algorithm. In each iteration, the algorithm first creates a set of searching directions, named “patterns”, which *positively spans* the entire feasible set. It then searches the directions one by one in order to find a direction, along which the objective value decreases. In each search, the algorithm needs to evaluate the objective value of $\mathbf{EP}_j(\alpha)$ given a GLB decision α , which can be obtained by plugging the optimal solution $q_j^*(p; \alpha)$ into the objective function of $\mathbf{EP}_j(\alpha)$. In this manner, our adapted GPS algorithm works like gradient-based algorithms, but without the need to compute gradient/subgradient. We summarize our proposed nested-loop algorithm in Algorithm 1.

In general, GPS algorithm is not guaranteed to converge to the globally optimal solution [20]. In the following theorem, we prove that our Algorithm 1 converges to the optimal solution to the convex problem **P3**, under proper conditions.

Theorem 3: Assume that $f_{U_j}(u)$, $j = 1, \dots, N$, are differentiable and their derivatives are continuous. Algorithm 1

converges to a globally optimal solution to **P3**, which is also an optimal solution to **P1** and **P2**.

V. COMPLEXITY AND PRACTICAL CONSIDERATIONS

In this section, we discuss the computation complexity and some practical considerations for our solution.

A. Computational Complexity

In our model and analysis, we assume that both MCP P_j and the demand U_j are continuous random variables. When applying them to practice, we need to sample a PDF (which is a continuous function) into a probability mass function (which is a discrete sequence). So we assume that we sample both the PDF of P_j , i.e., $f_{P_j}(p)$, and the PDF of U_j , i.e., $f_{U_j}(v)$, into sequences with length m . The value of m depends on both the ranges of MCP and demand and the accuracy we aim to achieve. Based on such sampling, we show the computational complexity of our proposed solution, i.e., Algorithm 1.

Theorem 4: If Algorithm 1 converges in n_{iter} iterations, its time complexity is $O(n_{\text{iter}}((N^5 m \log(Nm) + N^3 m^2)))$.

The complexity is linear with the number of iterations until convergence. However, exactly characterizing the convergence rate of GPS algorithm is still an open problem [9], and thus it is hard to get sharp bounds for the number of iterations, i.e., n_{iter} . Instead, we empirically evaluate the convergence rate of our Algorithm 1 in Sec. VI-C. The results show that our Algorithm 1 converges fast – within 30 iterations – for the practical setting considered (i.e., $n_{\text{iter}} \leq 30$). Considering that N , i.e., the number of datacenters of one CSP, is not very large, Theorem 4 shows that the complexity of our Algorithm 1 is affordable in practice.

Due to the space limitation, the discussions on the impact of imperfect probability distributions and only allowing limited number of bids are included in our technical report [38].

VI. EMPIRICAL EVALUATIONS

In this section, we use trace-driven simulations to empirically evaluate the performance of our proposed solution.

A. Dataset and Settings

Network Settings. We consider a CSP operating 3 datacenters in San Diego, Houston, and New York City, respectively. We assume that due to quality of experience consideration, the CSP cannot balance workloads between datacenters in San Diego and New York City. We set the unit bandwidth cost of routing workloads across datacenters as $z_{ij} = \kappa \cdot (\mu_1^{\text{RT}} + \mu_2^{\text{RT}} + \mu_3^{\text{RT}}) / 3$ if $i \neq j$, and $z_{ii} = 0$, $i = 1, 2, 3$. We let $\kappa = 0.1$ as a default setting, and we vary the values of κ to evaluate the overall cost-saving performance under different bandwidth-cost settings.

Workload and Electricity Demand. We get the numbers of service requests per hour against the Akamai CDN in North America for 48 days from Akamai's Internet Observatory website [1]. By using the conversion ratio claimed by Google for its datacenters [27], we scale up the request information to create an electricity demand series with averaged hourly demand of 125 MWh. The total demand is divided into three regions according to regional electricity consumptions of the three locations [38]. We set the ratio of demand of region i to be served locally, i.e., λ_i , to be 0.7. We also set datacenter j 's




Fig. 6. Empirical distributions of MCPs, 2pm.




Fig. 7. Empirical distributions of electricity demands, 2pm.

capacity C_j to be 30% larger than region j 's peak demand, since it was reported that on average 30% or more of the capacity of datacenters is idling in operation [6], [11].

Electricity Prices in Day-ahead and Real-time Markets.

We obtain the electricity prices (MCP of day-ahead market and real-time market price) from three regional ISO websites which serve the customers in San Diego, Houston, and New York, respectively [5] [10] [25]. The discounting factor β of selling back unused electricity is set as 0.5, which means that the CSP suffers half loss in case of over-supply.

Evaluation and Comparison. We test our design on 24 instances, each corresponding to one hour of the day. For each hour, the distributions of electricity demand, day-ahead MCP and real-time prices are learned from our dataset, and the real-time price expectation is computed from the distribution accordingly. For illustration purpose, we plot the empirical distributions of MCPs and demands for 2pm in Fig. 6 and Fig. 7, respectively. We denote our solution as **OptBidding-OptGLB**, in which the GPS part is based on an implementation used in [7], [8]. We test the following four baseline alternatives. (i) **NoBidding-NoGLB**: it represents the strategy of buying all electricity in real-time markets without doing GLB. It serves as the *benchmark* to compute cost reduction for other algorithms. (ii) **OptBidding-NoGLB**: it represents the strategy of optimally bidding in day-ahead markets but without doing GLB. (iii) **NoBidding-OptGLB**: it represents the strategy of doing no bidding in the day-ahead markets but purchasing all electricity in real-time markets and doing optimal GLB (adapted from the solution in [30]). (iv) **SimpleBidding-OptGLB**: it represents a joint bidding and GLB strategy proposed in [6], in which the CSP only submits one bid to each day-ahead market j with bidding price being μ_j^{RT} and the bidding quantity being the expectation of the GLB-allocated demand, and the CSP performs optimal GLB, accordingly.

B. Performance Comparison and Impact of Finite Bids

We compare the performance of different solutions in terms of the expected daily cost in Tab. I. Further, we also evaluate the performance loss due to that we approximate the optimal bidding curve (which may require the CSP to submit infinite number of bids) by using only 1 and 3 bids in our solution. We show the cost reduction of using infinite number of bids, 1 bid, and 3 bids in the last three rows of Tab. I, respectively.

We have the following observations. First of all, as seen from Tab. I, we can see that our proposed solution outperforms all other alternatives and reduces the CSP's operating cost by 20.8% as compared to the benchmark **NoBidding-NoGLB**. Meanwhile, we observe that **SimpleBidding-OptGLB** only reduces the cost by 3.8%, which is much less than that

TABLE I
Cost-saving performance of different schemes.

Solution	Daily Cost (k\$)	Reduction (%)
NoBidding-NoGLB	161.9	-
NoBidding-OptGLB (adapted from [30])	154.5	4.6
SimpleBidding-OptGLB [6]	155.8	3.8
OptBidding-NoGLB	135.4	16.4
OptBidding-OptGLB (Our solution)	128.2	20.8
OptBidding-OptGLB (1 bid)	133.3	17.7
OptBidding-OptGLB (3 bids)	128.6	20.5




Fig. 8. Optimal bidding curves for three day-ahead markets, 4pm.




Fig. 9. Objective values in each iteration of our Algorithm 1.

achieved by our solution OptBidding-OptGLB. Moreover, the cost reduction (3.8%) is even less than NoBidding-OptGLB (4.6%), which does not perform bidding in the day-ahead markets but purchases all electricity from the real-time markets. This highlights the importance of designing intelligent strategies for bidding on the day-ahead markets.

In addition to intelligent bidding strategy design, we observe that GLB also brings extra cost saving for CSP. For example, NoBidding-OptGLB reduces the cost by 4.6% as compared to NoBidding-NoGLB, and OptBidding-OptGLB achieves 4.4% extra reduction as compared to OptBidding-NoGLB.

Here, we use a simple method explained in our technical report [38] to approximate the optimal bidding curve with a finite number of bids (in particular, 1 and 3 bids in this experiment). From the last two rows in Tab. I, we observe that submitting 1 bid can achieve reasonably good performance (17.7% vs 20.8%). Submitting 3 bids can almost achieve the same performance as submitting infinite number of bids (20.5% vs 20.8%). This observation suggests that our solution performs well in practice even if the CSP is only allowed to submit a small number of bids to a day-ahead market. To understand this observation, we visualize the optimal bidding curves of three datacenters for one optimization instance (4pm) in Fig. 8. We can see that all three bidding curves are “flat” and thus can be accurately approximated by step-wise functions corresponding to submitting only a small number of bids.

C. Convergence Rate of the Joint Bidding and GLB Algorithm

In this subsection, we empirically evaluate the convergence rate of our proposed Algorithm 1. We run our algorithm for two instances with workload/price distribution of 10am and 2pm, respectively. From Fig. 9, we can see that our algorithm converges rather fast – within 30 iterations.

D. Impact of Demand Uncertainty / Distribution Estimation

To study the impact of demand uncertainty, we properly scale the electricity demand of all three regions such that the demand expectations stay the same and the average of the




Fig. 10. Cost reductions with different levels of demand uncertainty and different estimated distributions.




Fig. 11. Cost reductions when more workloads must be locally served, under different bandwidth cost.

normalized sample standard deviations among all three regions changes from 0.02 to 0.13, to mimic low to high uncertainty in workload demand. Here normalized sample standard deviation is defined as the ratio of the sample standard deviation to the sample mean. We apply our solution OptBidding-OptGLB to the set of scaled demands and plot the cost reduction in From Fig. 10. From Fig. 10, we can see that the cost reduction decreases as the demand certainty increases, but the performance loss is minor, suggesting that our solution OptBidding-OptGLB is robust to demand uncertainty.

We also study the impact of distribution estimation. In our solution OptBidding-OptGLB, we use the distribution of the demand U_j for region j as input. In practice, however, the CSP may not have the exact demand distributions, but just their estimates based on historical data. It is common for these estimated distributions to have the same mean and variance as the actual demand distributions, but it is difficult, if not impossible, for the estimated distribution to match the actual distribution exactly. A central question is then how sensitive is the performance of our solution OptBidding-OptGLB to the accuracy of the distribution estimation, given that we have obtained an accurate estimate of the mean and variance?

We explore answers to this question by comparing the performance achieved by our solution OptBidding-OptGLB based on the following distributions for demand with the same mean and variance: actual distribution, *normal distribution*, and *uniform distribution*. We compare their cost reductions in Fig. 10. As seen, the performance loss is minor, implying that accurate first and second order statistics of the demand distribution may be enough to determine the performance of our solution OptBidding-OptGLB. This observation also suggests an interesting direction for future work.

E. Impact of Local Service Requirement and Bandwidth Cost

We investigate the impact of local service requirement, where we changes the percentage of demand that must be served locally, i.e., λ_i , from 0.5 to 1.0. The simulation results are in Fig. 11, where we can see the cost reduction of our solution OptBidding-OptGLB decreases as λ_i increases. This matches our intuition that larger λ_i means that the CSP has less room to do GLB. When $\lambda_i = 1$, i.e., all demand should be served locally, our solution OptBidding-OptGLB coincides with OptBidding-NoGLB.

We also study the impact of bandwidth cost, where we choose two different values (0.1 and 0.4) for the bandwidth cost factor κ . We show the cost reduction in Fig. 11. As seen, a larger κ , meaning higher bandwidth cost, leads to smaller reduction, which matches out intuition.

VII. RELATED WORK

The seminal works [27], [28] propose the idea of GLB to effectively reduce electricity cost of datacenter operators. Later on many works [21], [23], [30], [35] have broadened the landscape of GLB with more practical considerations and design spaces. Instead of studying the benefit of GLB, several works [6], [33] study the impact of GLB on the electricity supply chain and electricity markets. Several works investigate how GLB should be operated in the presence of demand uncertainty and/or market price uncertainty. Both [29] and [36] utilize the long-term forward contracts to reduce operation risk. Similar to our work, [6], [13], [14], [34] deal with uncertainty by bidding in the day-ahead market. However, [13], [14] do not fully exploit the bidding design space, and [6], [34] do not consider demand uncertainty. Instead, our work fully exploits the bidding design space and simultaneously considers the demand and market uncertainty. Note that the problem of designing the optimal bids/offers has also been considered in [2], [3], [12], [22], [26] *etc.*, under various pricing models, but for a single-regional-market scenario. In particular, the authors in [2] design the optimal offer strategies for renewable generation company with given day-ahead market prices but uncertain wind generation.

VIII. CONCLUDING REMARKS

We develop an algorithm that is proven to minimize the total electricity and bandwidth cost of a CSP in face of workload and price uncertainty, by jointly optimizing strategic bidding in wholesale markets and GLB. Evaluations based on real-world traces show that our algorithm can reduce the CSP's cost by up to 20%. We show that, interestingly, while variability in workloads deteriorates cost saving, variability in market prices allows us to achieve a cost reduction even *larger* than the setting without price variability. Our work has several limitations. First, we assume that all the day-ahead market's MCPs and real-time market prices are mutually independent, which may not hold in practice. Second, we assume that the CSP has negligible market power, which may not hold for local markets even though globally datacenters today only consume less than 2% of the total electricity [11]. Finally, we do not model possible strategic behaviours of other market participants. Addressing these limitations is an interesting future direction.

REFERENCES

- [1] "Akamai Internet observatory," available at <https://www.akamai.com>
- [2] E. Y. Bitar, R. Rajagopal, P. P. Khargonekar, K. Poolla, and P. Varaiya, "Bringing wind energy to market," *IEEE Trans. Power Syst.*, vol. 27, no. 3, pp. 1225–1235, 2012.
- [3] T. K. Boomsma, N. Juul, and S.-E. Fleten, "Bidding in sequential electricity markets: The Nordic case," *European Journal of Operational Research*, vol. 238, no. 3, pp. 797–809, 2014.
- [4] R. Brown, "Report to congress on server and data center energy efficiency: Public law 109-431," *Lawrence Berkeley National Lab*, 2008.
- [5] "CAISO archive," available at <http://www.caiso.com>.
- [6] J. Camacho, Y. Zhang, M. Chen, and D. M. Chiu, "Balance your bids before your bits: The economics of geographic load-balancing," in *Proc. ACM e-Energy*, 2014.
- [7] A. L. Custódio, H. Rocha, and L. N. Vicente, "Incorporating minimum frobenius norm models in direct search," *Computational Optimization and Applications*, vol. 46, no. 2, pp. 265–278, 2010.
- [8] A. L. Custódio and L. N. Vicente, "Using sampling and simplex derivatives in pattern search methods," *SIAM Journal on Optimization*, vol. 18, no. 2, pp. 537–555, 2007.
- [9] E. D. Dolan, R. M. Lewis, and V. Torczon, "On the local convergence of pattern search," *SIAM Journal on Optimization*, vol. 14, no. 2, pp. 567–583, 2003.
- [10] "ERCOT archive," available at <http://www.ercot.com>.
- [11] "Facts about data centers," available at <http://energy.gov>.
- [12] S.-E. Fleten and E. Pettersen, "Constructing bidding curves for a price-taking retailer in the Norwegian electricity market," *IEEE Trans. Power Syst.*, vol. 20, no. 2, pp. 701–708, 2005.
- [13] M. Ghamkhari, H. Mohsenian-Rad, and A. Wierman, "Optimal risk-aware power procurement for data centers in day-ahead and real-time electricity markets," in *Proc. INFOCOM Workshop on SDP*, 2014.
- [14] M. Ghamkhari, A. Wierman, and H. Mohsenian-Rad, "Energy portfolio optimization of data centers," *IEEE Trans. Smart Grid*, 2016.
- [15] "Google energy wiki," http://en.wikipedia.org/wiki/Google_Energy.
- [16] Y. Guo and Y. Fang, "Electricity cost saving strategy in data centers by using energy storage," *IEEE Trans. Parallel Distrib. Syst.*, vol. 24, no. 6, pp. 1149–1160, 2013.
- [17] R. Herranz, A. M. San Roque, J. Villar, and F. A. Campos, "Optimal demand-side bidding strategies in electricity spot markets," *IEEE Trans. Power Syst.*, vol. 27, no. 3, pp. 1204–1213, 2012.
- [18] M. Khouja, "The single-period (news-vendor) problem: Literature review and suggestions for future research," *Omega*, vol. 27, no. 5, pp. 537–553, 1999.
- [19] J. Koomey, "Growth in data center electricity use 2005 to 2010," *A report by Analytical Press, The New York Times*, 2011.
- [20] R. M. Lewis and V. Torczon, "Pattern search methods for linearly constrained minimization," *SIAM Journal on Optimization*, vol. 10, no. 3, pp. 917–941, 2000.
- [21] M. Lin, Z. Liu, A. Wierman, and L. L. Andrew, "Online algorithms for geographical load balancing," in *Proc. IGCC*, 2012.
- [22] G. Liu, Y. Xu, and K. Tomovic, "Bidding strategy for microgrid in day-ahead market based on hybrid stochastic/robust optimization," *IEEE Trans. Smart Grid*, vol. 7, no. 1, pp. 227–237, 2016.
- [23] Z. Liu, M. Lin, A. Wierman, S. H. Low, and L. L. Andrew, "Greening geographical load balancing," in *Proc. ACM SIGMETRICS*, 2011.
- [24] B. Neupane, T. B. Pedersen, and B. Thiesson, "Evaluating the value of flexibility in energy regulation markets," in *Proc. ACM e-Energy*, 2015.
- [25] "NYISO archive," available at <http://www.nyiso.com>.
- [26] F. Paganini, P. Belzarena, and P. Monzón, "Decision making in forward power markets with supply and demand uncertainty," in *Proc. CISS*, 2014.
- [27] A. Qureshi, R. Weber, H. Balakrishnan, J. Gutttag, and B. Maggs, "Cutting the electric bill for Internet-scale systems," in *Proc. ACM SIGCOMM*, 2009.
- [28] L. Rao, X. Liu, and W. Liu, "Minimizing electricity cost: Optimization of distributed Internet data centers in a multi-electricity-market environment," in *Proc. IEEE INFOCOM*, 2010.
- [29] L. Rao, X. Liu, L. Xie, and Z. Pang, "Hedging against uncertainty: A tale of Internet data center operations under smart grid environment," *IEEE Trans. Smart Grid*, vol. 2, no. 3, pp. 555–563, 2011.
- [30] H. Shao, L. Rao, Z. Wang, X. Liu, Z. Wang, and K. Ren, "Optimal load balancing and energy cost management for Internet data centers in deregulated electricity markets," *IEEE Trans. Parallel and Distrib. Syst.*, vol. 25, no. 10, pp. 2659–2669, 2014.
- [31] Y. Shi, B. Xu, B. Zhang, and D. Wang, "Leveraging energy storage to optimize data center electricity cost in emerging power markets," in *Proc. ACM E-Energy*, 2016.
- [32] Q. Sun, S. Ren, C. Wu, and Z. Li, "An online incentive mechanism for emergency demand response in geo-distributed colocation data centers," in *Proc. ACM e-Energy*, 2016.
- [33] P. Wang, L. Rao, X. Liu, and Y. Qi, "D-pro: Dynamic data center operations with demand-responsive electricity prices in smart grid," *IEEE Trans. Smart Grid*, vol. 3, no. 4, pp. 1743–1754, 2012.
- [34] P. Wang, Y. Zhang, L. Deng, M. Chen, and X. Liu, "Second chance works out better: Saving more for data center operator in open energy market," in *Proc. CISS*, 2016.
- [35] R. Wang, N. Kandasamy, C. Nwankpa, and D. R. Kaeli, "Datacenters as controllable load resources in the electricity market," in *Proc. IEEE ICDCS*, 2013.
- [36] L. Yu, T. Jiang, Y. Cao, and Q. Zhang, "Risk-constrained operation for Internet data centers in deregulated electricity markets," *IEEE Trans. Parallel Distrib. Syst.*, vol. 25, no. 5, pp. 1306–1316, 2014.
- [37] X.-P. Zhang, *Restructured Electric Power Systems: Analysis of Electricity Markets with Equilibrium Models*. John Wiley & Sons, 2010.
- [38] Y. Zhang, L. Deng, M. Chen, and P. Wang, "Joint bidding and geographical load balancing for datacenters: Is uncertainty a blessing or a curse?" *CUHK, Tech. Rep. 2016*, <https://staff.ie.cuhk.edu.hk/%7Emhchen/papers/BGLBInfocomTech.pdf>.